

# S3 Security Best Practices for Cleanit

## 🔒 The Correct Architecture

```
User Request → Backend API → Private S3 Bucket → CloudFront CDN → User  
↓  
Presigned URLs / CloudFront URLs
```

## ⚠️ Never Do This (Public Bucket)

```
javascript  
  
// ❌ WRONG - DON'T DO THIS  
const params = {  
  Bucket: BUCKET_NAME,  
  Key: filePath,  
  Body: file.buffer,  
  ACL: 'public-read' // ❌ Makes file publicly accessible  
};
```

## ✓ Always Do This (Private Bucket)

```
javascript  
  
// ✅ CORRECT - Private bucket  
const params = {  
  Bucket: BUCKET_NAME,  
  Key: filePath,  
  Body: file.buffer,  
  ServerSideEncryption: 'AES256' // ✅ Private + encrypted  
};
```

## 🎯 Implementation Options

### Option 1: Private S3 + CloudFront (RECOMMENDED)

**Best for:** Production apps, scalable solutions, cost efficiency

**Advantages:**

- ✅ Secure (bucket is private)
- ✅ Fast (CloudFront CDN caching)
- ✅ Cheaper at scale (CloudFront pricing < S3 pricing)

- DDoS protection
- HTTPS by default

### Setup:

1. Create private S3 bucket
2. Create CloudFront distribution
3. Set Origin Access Identity (OAI) so only CloudFront can access S3
4. Backend returns CloudFront URLs

### Implementation:

```
javascript

// Upload returns CloudFront URL
const imageUrl = await uploadToS3(file);
// Returns: https://d123456.cloudfront.net/uploads/2024/01/abc123.jpg

// Store this URL in database - it's already accessible through CloudFront
await prisma.campaign.create({
  data: {
    imageUrl: imageUrl // CloudFront URL
  }
});
```

**Terraform setup:** See updated [main.tf](#) artifact

---

### Option 2: Private S3 + Presigned URLs

**Best for:** MVP, simpler setup, private content

#### Advantages:

- Secure (bucket is private)
- Time-limited access
- No CloudFront setup needed
- Good for MVP

#### Disadvantages:

-  Higher costs at scale (S3 data transfer pricing)
-  No caching (slower for users)

- ⚠ URLs expire (need regeneration)

## Implementation:

```
javascript

// Upload returns S3 key (not full URL)
const s3Key = await uploadToS3(file);
// Returns: uploads/2024/01/abc123.jpg

// Store S3 key in database
await prisma.campaign.create({
  data: {
    imageUrl: s3Key // Just the key, not full URL
  }
});

// When serving to frontend, generate presigned URL
const campaign = await prisma.campaign.findUnique({ where: { id } });
campaign.imageUrl = await generatePresignedUrl(campaign.imageUrl, 3600); // 1 hour
```

## Modified Campaign Controller:

```
javascript

exports.getCampaignById = async (req, res, next) => {
  try {
    const campaign = await prisma.campaign.findUnique({
      where: { id: req.params.id }
    });

    // Generate presigned URL for the image
    campaign.imageUrl = await storageService.generatePresignedUrl(
      campaign.imageUrl,
      3600 // URL valid for 1 hour
    );

    res.json({ campaign });
  } catch (error) {
    next(error);
  }
};
```

## Option 3: Backend Proxy (Most Secure, Complex)

**Best for:** Highly sensitive content, full access control

## Advantages:

- Complete control over who accesses what
- Can add authentication, logging, watermarks
- No presigned URL expiration issues

## Disadvantages:

- Backend handles all traffic (expensive, slower)
- More complex to implement
- Backend becomes bottleneck

## Implementation:

javascript

```
// New route in backend
router.get('/images/:key', authenticateToken, async (req, res) => {
  const s3Stream = s3.getObject({
    Bucket: BUCKET_NAME,
    Key: req.params.key
  }).createReadStream();

  s3Stream.pipe(res);
});

// Frontend uses: /api/images/uploads/2024/01/abc123.jpg
```

## 📊 Comparison Table

Feature	Public S3	CloudFront + Private S3	Presigned URLs	Backend Proxy
Security	<input checked="" type="checkbox"/> Poor	<input checked="" type="checkbox"/> Excellent	<input checked="" type="checkbox"/> Excellent	<input checked="" type="checkbox"/> Excellent
Speed	<input checked="" type="checkbox"/> OK	<input checked="" type="checkbox"/> Fast	<input checked="" type="checkbox"/> OK	<input checked="" type="checkbox"/> Slow
Cost (high traffic)	<input checked="" type="checkbox"/> High	<input checked="" type="checkbox"/> Low	<input checked="" type="checkbox"/> Medium	<input checked="" type="checkbox"/> Very High
Setup Complexity	<input checked="" type="checkbox"/> Simple	<input checked="" type="checkbox"/> Medium	<input checked="" type="checkbox"/> Simple	<input checked="" type="checkbox"/> Complex
Scalability	<input checked="" type="checkbox"/> Good	<input checked="" type="checkbox"/> Excellent	<input checked="" type="checkbox"/> Good	<input checked="" type="checkbox"/> Poor
Recommended	Never	Production	MVP	Rarely

# Recommended Approach for Cleanit

## Phase 1: MVP (Presigned URLs)

Start simple with presigned URLs:

- Private S3 bucket
- Generate presigned URLs on-demand
- Good enough for initial launch
- Easy to implement

## Phase 2: Production (CloudFront)

Scale with CloudFront:

- Add CloudFront distribution
- Point to private S3 bucket
- Update storage service to return CloudFront URLs
- Better performance + lower costs

---

## Implementation Steps

### Step 1: Create Private S3 Bucket

```
bash

aws s3api create-bucket \
--bucket cleanit-uploads \
--region ap-south-1 \
--create-bucket-configuration LocationConstraint=ap-south-1

# Block all public access
aws s3api put-public-access-block \
--bucket cleanit-uploads \
--public-access-block-configuration \
"BlockPublicAcls=true,IgnorePublicAcls=true,BlockPublicPolicy=true,RestrictPublicBuckets=true"

# Enable encryption
aws s3api put-bucket-encryption \
--bucket cleanit-uploads \
--server-side-encryption-configuration \
'{"Rules":[{"ApplyServerSideEncryptionByDefault":{"SSEAlgorithm":"AES256"}]}]'
```

## Step 2: Verify Bucket is Private

```
bash

# Should return: Public Access: Blocked
aws s3api get-public-access-block --bucket cleanit-uploads
```

## Step 3: Update Backend Code

Use the updated `storageService.js` provided above.

## Step 4: Test

```
javascript

// Upload test
const result = await uploadToS3(testFile);
console.log(result); // Should be S3 key or CloudFront URL

// Try accessing directly (should fail if private)
// https://cleanit-uploads.s3.ap-south-1.amazonaws.com/uploads/test.jpg
// Expected: Access Denied

// Generate presigned URL (should work)
const url = await generatePresignedUrl(result);
console.log(url); // Should work for 1 hour
```

## Checklist for Production

- S3 Block Public Access enabled
- S3 bucket encryption enabled
- S3 versioning enabled (for backup)
- CloudFront distribution created
- CloudFront uses HTTPS only
- Origin Access Identity configured
- S3 bucket policy allows only CloudFront
- Backend stores CloudFront URLs (not S3 URLs)
- Removed any `ACL: 'public-read'` from code
- Test that direct S3 URLs don't work
- Test that CloudFront URLs work
- Set up CloudWatch alarms for S3 access
- Enable S3 access logging
- Configure CORS properly

## Cost Comparison (Example: 1TB transfer/month)

Method	Approximate Cost
Public S3	~\$90/month
Private S3 + Presigned URLs	~\$90/month
CloudFront + Private S3	~\$60/month

**Winner:** CloudFront (33% cheaper + faster)

---

## Security Best Practices

1. **Never use public buckets** for user-uploaded content
  2. **Always encrypt** data at rest (use S3 encryption)
  3. **Use HTTPS only** (CloudFront enforces this)
  4. **Enable versioning** for backup/recovery
  5. **Monitor access** with CloudWatch and CloudTrail
  6. **Rotate credentials** regularly
  7. **Use IAM roles** instead of hardcoded keys when possible
  8. **Scan uploads** for malware (add to future roadmap)
  9. **Set size limits** (already done in multer config)
  10. **Validate file types** (already done in upload middleware)
- 

## Troubleshooting

"Access Denied" when accessing images

If using presigned URLs:

- Check URL hasn't expired
- Verify IAM permissions for backend
- Check bucket policy

If using CloudFront:

- Verify OAI is configured correctly
- Check S3 bucket policy allows CloudFront
- Wait 15-30 minutes for CloudFront distribution to deploy

## Images load slowly

- Enable CloudFront caching
- Set appropriate Cache-Control headers
- Use CloudFront over S3 direct access

## High AWS bills

- Check if bucket accidentally became public
  - Enable CloudFront (cheaper than S3 at scale)
  - Set lifecycle policies to delete old files
  - Monitor with AWS Cost Explorer
- 

## Additional Resources

- [AWS S3 Security Best Practices](#)
  - [CloudFront with S3](#)
  - [S3 Presigned URLs](#)
- 

## Summary

### For Cleanit MVP:

1. Use **PRIVATE S3 bucket** (never public)
2. Start with **presigned URLs** (simpler)
3. Migrate to **CloudFront** before production (better performance + cost)
4. Never set `ACL: 'public-read'` in your code

The updated code provided above implements this correctly! 